

Para hacer efectivos los cambios, reiniciamos el servicio con `sudo systemctl nginx restart`. No obstante, al intentar acceder al servidor nginx a través de su IP, `192.168.56.104`, nos sigue sirviendo su contenido. Para solucionar esto debemos eliminar la línea que configura nginx como servidor web en `/etc/nginx/nginx.conf`.

```
# gzip_proxied any;
# gzip_comp_level 6;
# gzip_buffers 16 8k;
# gzip_http_version 1.1;
# gzip_types text/plain text/css application/json application/javascript text/xml applicatio
n/xml application/xml+rss text/javascript;

##
# Virtual Host Configs
##

include /etc/nginx/conf.d/*.conf;
# include /etc/nginx/sites-enabled/*;
}

#mail {
#   # See sample authentication script at:
#   # http://wiki.nginx.org/ImapAuthenticateWithApachePhpScript
#
#   # auth_http localhost/auth.php;
#   # pop3_capabilities "TOP" "USER";
#   # imap_capabilities "IMAP4rev1" "UIDPLUS";
#
#   server {
#       listen     localhost:110;
#       protocol   pop3;
#       proxy      on;
#   }
#
#   server {
#       listen     localhost:143;
#       protocol   imap;
#       proxy      on;
#   }
#}
-- VISUAL --
```

38 62,39-46 Final

Ahora mismo, si accedemos desde un navegador o descargamos el contenido del servidor M3 mediante `curl`, debería servirnos el contenido de M1 y M2 por turnos, ya que ambas máquinas tienen el mismo peso.

Podemos observarlo aquí:

```
patchispatch@m1:~$ curl http://192.168.56.104
<!DOCTYPE html>
<html>
  <head>
    M1
  </head>
  <body>
    <h1>Este es el index de M1</h1>
  </body>
</html>
patchispatch@m1:~$ curl http://192.168.56.104
<!DOCTYPE html>
<html>
  <head>
    M2
  </head>
  <body>
    <h1>Este es el index de M2</h1>
  </body>
</html>
patchispatch@m1:~$ _
```

Configurar una máquina e instalar haproxy como balanceador de carga

Ya tenemos la máquina M3, por lo que instalamos haproxy con el comando `sudo apt install haproxy`. Para que todo funcione correctamente, desactivamos nginx con `sudo systemctl stop nginx`.

Para configurar haproxy como un balanceador de carga debemos indicar nuestros servidores M1 y M2, así como el tipo de peticiones a balancear, en el archivo `/etc/haproxy/haproxy.cfg`:

```
# Default SSL material locations
ca-base /etc/ssl/certs
crt-base /etc/ssl/private

# Default ciphers to use on SSL-enabled listening sockets.
# For more information, see ciphers(1SSL). This list is from:
# https://hynek.me/articles/hardening-your-web-servers-ssl-ciphers/
# An alternative list with additional directives can be obtained from
# https://mozilla.github.io/server-side-tls/ssl-config-generator/?server=haproxy
ssl-default-bind-ciphers ECDH+AESGCM:DH+AESGCM:ECDH+AES256:DH+AES256:ECDH+AES128:DH+AES:RSA+
AESGCM:RSA+AES:!aNULL:!MD5:!DSS
ssl-default-bind-options no-ssl3

defaults
    log      global
    mode     http
    option   httplog
    option   dontlognull
    timeout connect 5000
    timeout client  50000
    timeout server  50000
    errorfile 400 /etc/haproxy/errors/400.http
    errorfile 403 /etc/haproxy/errors/403.http
    errorfile 408 /etc/haproxy/errors/408.http
    errorfile 500 /etc/haproxy/errors/500.http
    errorfile 502 /etc/haproxy/errors/502.http
    errorfile 503 /etc/haproxy/errors/503.http
    errorfile 504 /etc/haproxy/errors/504.http

frontend http-in
    bind *:80
    default_backend servidoresSWAP

backend servidoresSWAP
    server m1 192.168.56.103:80 maxconn 32
    server m2 192.168.56.101:80 maxconn 32
"/etc/haproxy/haproxy.cfg" 45L, 1441C escritos
```

45,39-46 Final

Lanzamos haproxy con `sudo systemctl restart haproxy` y comprobamos el correcto balanceo de la carga:

```
patchispatch@m1:~$ curl http://192.168.56.104
<!DOCTYPE html>
<html>
  <head>
    M1
  </head>
  <body>
    <h1>Este es el index de M1</h1>
  </body>
</html>
patchispatch@m1:~$ curl http://192.168.56.104
<!DOCTYPE html>
<html>
  <head>
    M2
  </head>
  <body>
    <h1>Este es el index de M2</h1>
  </body>
</html>
patchispatch@m1:~$
```

Someter a la granja web a una alta carga, generada con la herramienta Apache Benchmark, teniendo primero nginx y después haproxy

Para probar la resistencia a la carga, ejecutamos el comando `ab -n 10000 -c 10 http://192.168.56.104/index.html`, tanto con nginx como con haproxy. Haremos la petición esta vez desde el host, por variar.

Prueba con nginx:

```

Benchmarking 192.168.56.104 (be patient)
Completed 1000 requests
Completed 2000 requests
Completed 3000 requests
Completed 4000 requests
Completed 5000 requests
Completed 6000 requests
Completed 7000 requests
Completed 8000 requests
Completed 9000 requests
Completed 10000 requests
Finished 10000 requests

Server Software:      nginx/1.14.0
Server Hostname:      192.168.56.104
Server Port:          80

Document Path:        /index.html
Document Length:       106 bytes

Concurrency Level:     10
Time taken for tests:   2.170 seconds
Complete requests:      10000
Failed requests:        0
Total transferred:      3750000 bytes
HTML transferred:       1060000 bytes
Requests per second:    4609.30 [#/sec] (mean)
Time per request:       2.170 [ms] (mean)
Time per request:       0.217 [ms] (mean, across all concurrent requests)
Transfer rate:          1687.98 [Kbytes/sec] received

Connection Times (ms)
      min    mean[+/-sd] median    max
Connect:    0       0   0.1      0      4
Processing:  1       2   0.6      2      9
Waiting:    1       2   0.6      2      9
Total:      1       2   0.6      2      9

Percentage of the requests served within a certain time (ms)
 50%    2
 66%    2
 75%    2
 80%    2
 90%    3
 95%    3
 98%    4
 99%    5
100%    9 (longest request)

~ took 2s
>

```

Prueba con haproxy:

```

Benchmarking 192.168.56.104 (be patient)
Completed 1000 requests
Completed 2000 requests
Completed 3000 requests
Completed 4000 requests
Completed 5000 requests
Completed 6000 requests
Completed 7000 requests
Completed 8000 requests
Completed 9000 requests
Completed 10000 requests
Finished 10000 requests

Server Software:      Apache/2.4.29
Server Hostname:      192.168.56.104
Server Port:          80

Document Path:        /index.html
Document Length:      106 bytes

Concurrency Level:    10
Time taken for tests:  2.412 seconds
Complete requests:    10000
Failed requests:      0
Total transferred:    3760000 bytes
HTML transferred:     1060000 bytes
Requests per second:  4145.98 [#/sec] (mean)
Time per request:     2.412 [ms] (mean)
Time per request:     0.241 [ms] (mean, across all concurrent requests)
Transfer rate:        1522.35 [Kbytes/sec] received

Connection Times (ms)
              min    mean[+/-sd] median    max
Connect:      0      0   0.2      0      4
Processing:   1      2   0.6      2      9
Waiting:      1      2   0.6      2      9
Total:        1      2   0.6      2      9

Percentage of the requests served within a certain time (ms)
 50%      2
 66%      2
 75%      2
 80%      3
 90%      3
 95%      3
 98%      4
 99%      6
100%      9 (longest request)

~ took 2s
> _

```

Podemos ver que todas las peticiones se han completado con ambos balanceadores, y que los tests se han ejecutado ligeramente más rápido utilizando nginx.